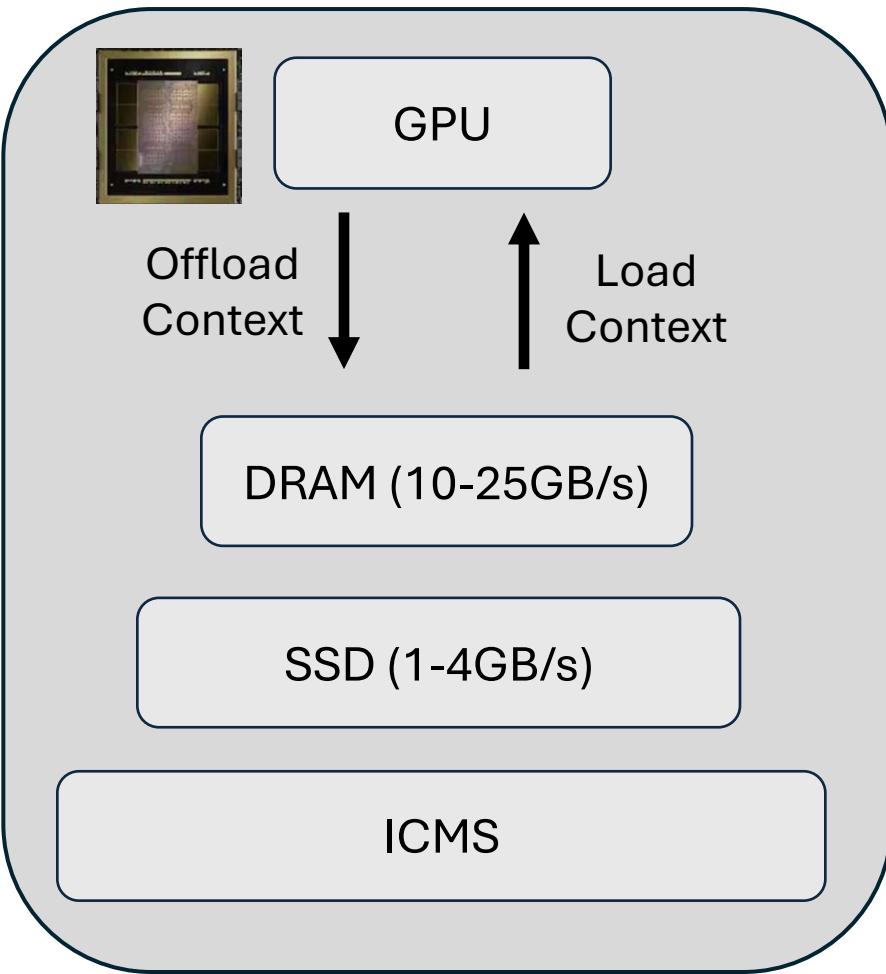


## Prefill Node (Compute Optimized)



## Decode Node (Memory Optimized)

